

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 class TwoLayerNet(object):
5     """
6     A two-layer fully-connected neural network. The net has an input
7     dimension of
8     N, a hidden layer dimension of H, and performs classification over C
9     classes.
10    We train the network with a softmax loss function and L2 regularization
11    on the
12    weight matrices. The network uses a ReLU nonlinearity after the first
13    fully
14    connected layer.
15
16    In other words, the network has the following architecture:
17
18    input - fully connected layer - ReLU - fully connected layer - MSE Loss
19
20    ReLU function:
21    (i)  $x = x$  if  $x \geq 0$  (ii)  $x = 0$  if  $x < 0$ 
22
23    The outputs of the second fully-connected layer are the scores for each
24    class.
25    """
26
27    def __init__(self, input_size, hidden_size, output_size, std=1e-4):
28        """
29        Initialize the model. Weights are initialized to small random values
30        and
31        biases are initialized to zero. Weights and biases are stored in the
32        variable self.params, which is a dictionary with the following keys:
33
34        W1: First layer weights; has shape (H, D)
35        b1: First layer biases; has shape (H,)
36        W2: Second layer weights; has shape (C, H)
37        b2: Second layer biases; has shape (C,)
38
39        Inputs:
40        - input_size: The dimension D of the input data.
41        - hidden_size: The number of neurons H in the hidden layer.
42        - output_size: The number of classes C.
43        """
44        self.params = {}
45        self.params['W1'] = std * np.random.randn(hidden_size, input_size)
46        self.params['b1'] = np.zeros(hidden_size)
47        self.params['W2'] = std * np.random.randn(output_size, hidden_size)
48        self.params['b2'] = np.zeros(output_size)
49
50    def loss(self, X, y=None, reg=0.0):
51        """
52        Compute the loss and gradients for a two layer fully connected neural
53        network.
54
55        Inputs:
56        - X: Input data of shape (N, D). Each X[i] is a training sample.
57        - y: Vector of training labels. y[i] is the label for X[i], and each
58        y[i] is
59        an integer in the range  $0 \leq y[i] < C$ . This parameter is optional;
60        if it

```

```

53         is not passed then we only return scores, and if it is passed then
we
54         instead return the loss and gradients.
55         - reg: Regularization strength.
56
57     Returns:
58     If y is None, return a matrix scores of shape (N, C) where scores[i,
c] is
59     the score for class c on input X[i].
60
61     If y is not None, instead return a tuple of:
62     - loss: Loss (data loss and regularization loss) for this batch of
training
63     samples.
64     - grads: Dictionary mapping parameter names to gradients of those
parameters
65     with respect to the loss function; has the same keys as
self.params.
66     """
67     # Unpack variables from the params dictionary
68     W1, b1 = self.params['W1'], self.params['b1']
69     W2, b2 = self.params['W2'], self.params['b2']
70     N, D = X.shape
71
72     # Compute the forward pass
73     scores = None
74
75     # ===== #
76     # START YOUR CODE HERE
77     # ===== #
78     # Calculate the output scores of the neural network. The result
79     # should be (N, C). As stated in the description for this class,
80     # there should not be a ReLU layer after the second fully-connected
81     # layer.
82     # The code is partially given
83     # The output of the second fully connected layer is the output
scores.
84     # Do not use a for loop in your implementation.
85     # Please use 'h1' as input of hidden layers, and 'a2' as output of
86     # hidden layers after ReLU activation function.
87     # [Input X] --W1,b1--> [h1] -ReLU-> [a2] --W2,b2--> [scores]
88     # You may simply use np.maximum for implementing ReLU.
89     # Note that there is only one ReLU layer.
90     # Note that please do not change the variable names (h1, h2, a2)
91     # ===== #
92
93     h1 = np.dot(X, W1.T) + b1
94     a2 = np.zeros(h1.shape)
95     a2 = np.maximum(a2, h1) # activation with input of h1
96     h2 = np.dot(a2, W2.T) + b2
97     scores = h2
98
99     # ===== #
100    # END YOUR CODE HERE
101    # ===== #
102
103
104    # If the targets are not given then jump out, we're done
105    if y is None:
106        return scores

```

```

107
108     # Compute the loss
109     loss = None
110
111     # scores is num_examples by num_classes (N, C)
112     def softmax_loss(x, y):
113         loss, dx = 0,0
114         #
===== #
115         # START YOUR CODE HERE (BONUS QUESTION)
116         #
===== #
117         # Calculate the cross entropy loss after softmax output layer.
118         # The format are provided in the notebook.
119         # This function should return loss and dx, same as MSE loss
function.
120         #
===== #
121
122         pass
123
124         #
===== #
125         # END YOUR CODE HERE
126         #
===== #
127         return loss, dx
128
129
130     def MSE_loss(x, y):
131         loss, dx = 0,0
132         #
===== #
133         # START YOUR CODE HERE
134         #
===== #
135         # This function should return loss and dx (gradients ready for
back prop).
136         # The loss is MSE loss between network ouput and one hot vector
of class
137         # labels is required for backpropagation.
138         #
===== #
139         # Hint: Check the type and shape of x and y.
140         # e.g. print('DEBUG:x.shape, y.shape', x.shape, y.shape)
141
142         # x is our y_pred, and y is the y_target
143
144         num_samples = x.shape[0]
145         num_attrs = x.shape[1]
146         y_target = np.zeros((num_samples, num_attrs))
147         for i in range(num_samples):
148             y_target[i][y[i]] = 1
149         diff = x - y_target
150         dx = diff / num_samples
151         loss = 0.5 * np.sum(np.square(diff)) / num_samples
152
153         #
===== #
154         # END YOUR CODE HERE

```

```

155         #
===== #
156         return loss, dx
157
158         # data_loss, dscore = softmax_loss(scores, y)
159         # The above line is for bonus question. If you have implemented
softmax_loss, de-comment this line instead of MSE error.
160
161         data_loss, dscore = MSE_loss(scores, y) # "comment" this line if you
use softmax_loss
162         # ===== #
163         # START YOUR CODE HERE
164         # ===== #
165         # Calculate the regularization loss. Multiply the regularization
166         # loss by 0.5 (in addition to the factor reg).
167         # ===== #
168         reg_loss = 0.5 * reg * (np.sum(W1*W1) + np.sum(W2*W2))
169
170         # ===== #
171         # END YOUR CODE HERE
172         # ===== #
173         loss = data_loss + reg_loss
174
175         grads = {}
176
177         # ===== #
178         # START YOUR CODE HERE
179         # ===== #
180         # Backpropagation: (You do not need to change this!)
181         # Backward pass is implemented. From the dscore error, we calculate
182         # the gradient and store as grads['W1'], etc.
183         # ===== #
184         grads['W2'] = a2.T.dot(dscore).T + reg * W2
185         grads['b2'] = np.ones(N).dot(dscore)
186
187         da_h = np.zeros(h1.shape)
188         da_h[h1>0] = 1
189         dh = (dscore.dot(W2) * da_h)
190
191         grads['W1'] = np.dot(dh.T,X) + reg * W1
192         grads['b1'] = np.ones(N).dot(dh)
193         # ===== #
194         # END YOUR CODE HERE
195         # ===== #
196
197         return loss, grads
198
199     def train(self, X, y, X_val, y_val,
200               learning_rate=1e-3, learning_rate_decay=0.95,
201               reg=1e-5, num_iters=100,
202               batch_size=200, verbose=False):
203         """
204         Train this neural network using stochastic gradient descent.
205
206         Inputs:
207         - X: A numpy array of shape (N, D) giving training data.
208         - y: A numpy array of shape (N,) giving training labels; y[i] = c
means that
209             X[i] has label c, where 0 <= c < C.
210         - X_val: A numpy array of shape (N_val, D) giving validation data.

```

```

211     - y_val: A numpy array of shape (N_val,) giving validation labels.
212     - learning_rate: Scalar giving learning rate for optimization.
213     - learning_rate_decay: Scalar giving factor used to decay the
learning rate
214         after each epoch.
215     - reg: Scalar giving regularization strength.
216     - num_iters: Number of steps to take when optimizing.
217     - batch_size: Number of training examples to use per step.
218     - verbose: boolean; if true print progress during optimization.
219     """
220     num_train = X.shape[0]
221     iterations_per_epoch = max(num_train / batch_size, 1)
222
223     # Use SGD to optimize the parameters in self.model
224     loss_history = []
225     train_acc_history = []
226     val_acc_history = []
227
228     for it in np.arange(num_iters):
229         X_batch = None
230         y_batch = None
231
232         # Create a minibatch (X_batch, y_batch) by sampling batch_size
233         # samples randomly.
234
235         b_index = np.random.choice(num_train, batch_size)
236         X_batch = X[b_index]
237         y_batch = y[b_index]
238
239         # Compute loss and gradients using the current minibatch
240         loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
241         loss_history.append(loss)
242
243         #
===== #
244         # START YOUR CODE HERE
245         #
===== #
246         # Perform a gradient descent step using the minibatch to update
247         # all parameters (i.e., W1, W2, b1, and b2).
248         # The gradient has been calculated as grads['W1'], grads['W2'],
249         # grads['b1'], grads['b2']
250         # For example,
251         # W1(new) = W1(old) - learning_rate * grads['W1']
252         # (this is not the exact code you use!)
253         #
===== #
254
255         self.params['W1'] = self.params['W1'] - learning_rate *
grads['W1']
256         self.params['b1'] = self.params['b1'] - learning_rate *
grads['b1']
257         self.params['W2'] = self.params['W2'] - learning_rate *
grads['W2']
258         self.params['b2'] = self.params['b2'] - learning_rate *
grads['b2']
259
260         #
===== #
261         # END YOUR CODE HERE

```

```

262         #
===== #
263
264         if verbose and it % 100 == 0:
265             print('iteration {} / {}: loss {}'.format(it, num_iters,
loss))
266
267         # Every epoch, check train and val accuracy and decay learning
rate.
268         if it % iterations_per_epoch == 0:
269             # Check accuracy
270             train_acc = (self.predict(X_batch) == y_batch).mean()
271             val_acc = (self.predict(X_val) == y_val).mean()
272             train_acc_history.append(train_acc)
273             val_acc_history.append(val_acc)
274
275             # Decay learning rate
276             learning_rate *= learning_rate_decay
277
278         return {
279             'loss_history': loss_history,
280             'train_acc_history': train_acc_history,
281             'val_acc_history': val_acc_history,
282         }
283
284     def predict(self, X):
285         """
286         Use the trained weights of this two-layer network to predict labels
for
287         data points. For each data point we predict scores for each of the C
288         classes, and assign each data point to the class with the highest
score.
289
290         Inputs:
291         - X: A numpy array of shape (N, D) giving N D-dimensional data points
to
292             classify.
293
294         Returns:
295         - y_pred: A numpy array of shape (N,) giving predicted labels for
each of
296             the elements of X. For all i, y_pred[i] = c means that X[i] is
predicted
297             to have class c, where 0 <= c < C.
298         """
299         y_pred = None
300
301         # ===== #
302         # START YOUR CODE HERE
303         # ===== #
304         # Predict the class given the input data.
305         # ===== #
306         scores = self.loss(X)
307         y_pred = np.argmax(scores, axis=1)
308
309         # ===== #
310         # END YOUR CODE HERE
311         # ===== #
312
313         return y_pred

```

314  
315  
316